


# Optimization – Quantization (optional)

- Quantization on Board


Computing Memory Architecture Lab.


# Quantization on Board


# New github repo

 [tahsd](#) / [hsd21\\_project\\_quant](#)

[Code](#) [Issues](#) [Pull requests](#) [Actions](#) [Projects](#) [Wiki](#) [Security](#) [Insights](#) [Settings](#)

 **main** ▾


 **1 branch**


 **0 tags**












[Go to file](#)

[Add file ▾](#)

[Code ▾](#)


 **tantara** **init**

253d6e0 2 days ago  **2 commits**

 <b>build</b>	init	2 days ago
 <b>data</b>	init	12 months ago
 <b>include</b>	init	2 days ago
 <b>pretrained_weights</b>	init	12 months ago
 <b>proto</b>	init	2 days ago
 <b>src</b>	init	2 days ago
 <b>Makefile</b>	init	12 months ago
 <b>benchmark.sh</b>	init	12 months ago
 <b>download.sh</b>	init	12 months ago
 <b>eval.py</b>	init	12 months ago
 <b>models.py</b>	init	12 months ago


# (pretrained) Convolutional Network Network(CNN)

- Input: **28x28** pixels  $\rightarrow$  **6 3x3 Conv**  $\rightarrow$  **30** values  $\rightarrow$  **10** values
  - 1<sup>st</sup> Conv: **28x28** inputs  $\rightarrow$  6 3x3 Conv  $\rightarrow$  6 26x26 outputs
  - 2<sup>nd</sup> FC:  $6 \cdot 26 \cdot 26 (=4056)$  inputs  $\rightarrow$  FC  $\rightarrow$  30 outputs **Lab 12(Quantization)**
  - 3<sup>rd</sup> FC: 30 inputs  $\rightarrow$  FC  $\rightarrow$  10 outputs **Lab 2(MV)**



# Review: Google's Int8 Solution

- Goal: running int8 neural networks on CPUs supporting int8 SIMD instructions
- Conversion, i.e., quantization from floating to 8-bit unsigned integer




# Review: Quantization Example

- Calculating scale (S) and zero (Z) for int8

$$\begin{array}{llll}
 -31.4 \rightarrow 0 & -31.4 = S \cdot (0 - Z) & SZ = 31.4 & Z = \text{ceil}(31.4 / 0.38) = 83 \\
 65.7 \rightarrow 255 & 65.7 = S \cdot (255 - Z) & 65.7 + 31.4 = 255S & S = 97.1 / 255 = 0.38
 \end{array}$$

- Example:  $10.5 \rightarrow \text{ceil}(10.5 / 0.38) + 83 = 29 + 83 = 112$



# Review: Integer Only Computation

- Matrix multiplication,  $r_3 = r_1 * r_2$ 
  - Weight matrix  $r_1$  and activation matrix  $r_2$

**Zero offset**

$$r_{\alpha}^{(i,j)} = S_{\alpha}(q_{\alpha}^{(i,j)} - Z_{\alpha})$$

What if  $Z_1$  and  $Z_2$  are zero?

$$S_3(q_3^{(i,k)} - Z_3) = \sum_{j=1}^N S_1(q_1^{(i,j)} - Z_1) S_2(q_2^{(j,k)} - Z_2)$$

$$q_3^{(i,k)} = Z_3 + M \sum_{j=1}^N (q_1^{(i,j)} - Z_1)(q_2^{(j,k)} - Z_2)$$

$$M := \frac{S_1 S_2}{S_3}$$

**Scaling(=dequantization)**

$$q_3^{(i,k)} = Z_3 + M \left( N Z_1 Z_2 - Z_1 a_2^{(k)} - Z_2 \bar{a}_1^{(i)} + \sum_{j=1}^N q_1^{(i,j)} q_2^{(j,k)} \right)$$

$$a_2^{(k)} := \sum_{j=1}^N q_2^{(j,k)}, \quad \bar{a}_1^{(i)} := \sum_{j=1}^N q_1^{(i,j)}$$


$$\sum_{j=1}^N q_1^{(i,j)} q_2^{(j,k)}$$

# Solution: 7b quantization + zero offset

- Calculating scale (S) and zero (Z) for 7b

$$\begin{array}{llll} -31.4 \rightarrow 0 & -31.4 = S \cdot (0 - Z) & SZ = 31.4 & Z = \text{ceil}(31.4 / 0.76) = 41 \\ 65.7 \rightarrow 127 & 65.7 = S \cdot (127 - Z) & 65.7 + 31.4 = 127S & S = 97.1 / 127 = 0.76 \end{array}$$

- Example:  $10.5 \rightarrow \text{ceil}(10.5 / 0.76) + 41 = 14 + 41 = 55$





# Install

---

- **Connect to the internet with LAN cable**
- Dependency download
  - \$ sudo apt-get update -y
  - \$ sudo apt-get install -y libprotobuf-dev protobuf-compiler python python-numpy
- Code download
  - \$ git clone [https://github.com/tahsd/hsd21\\_project\\_quant.git](https://github.com/tahsd/hsd21_project_quant.git)
- Dataset download
  - \$ bash download.sh

# Run

---

- Edit `src/fpga\_api\_on\_cpu.cpp`
  - void quantize(...), dequantize(...)
  - Also edit `src/fpga\_api.cpp` for optimization
- Build
  - \$ make
- Evaluate
  - \$ sudo python eval.py --num\_test\_images 100 --m\_size 8 --v\_size 8 -  
-network cnn --run\_type cpu **--quantized**
- Options
  - num\_test\_images : 1~10000
  - m\_size, v\_size : 8
  - network: cnn
  - run\_type: cpu(tutorial), fpga(optimized version)
    - fpga uses the IP that you implemented
  - **quantized: enable(8-bit) or disable(32-bit, full precision)**

# Implement Quantization - CPU

- Edit `src/fpga\_api\_on\_cpu.cpp`

```
const float* FPGA::blockMM(Compute* comp)
{
    num_block_call_ += 1;

    // cpu version
    int* m1 = this->qmatrix_M1();
    int* m2 = this->qmatrix_M2();
    float* out = reinterpret_cast<float*>(output_M);

    if(comp->quantized)
    {
        char act_bits_min = 0;
        char act_bits_max = (1<<(comp->act_bits-1))-1;

        float act_scale = 0; // TODO calculate the scale factor
        char act_offset = 0; // TODO calculate the zero-offset
        quantize(); // TODO complete quantize function

        char weight_bits_min = 0;
        char weight_bits_max = (1<<(comp->weight_bits-1))-1;

        float weight_scale = 0; // TODO calculate the scale factor
        char weight_offset = 0; // TODO calculate the zero-offset
        quantize(); // TODO complete quantize function

        for(int i = 0; i < v_size_; ++i)
        {
            for(int j = 0; j < v_size_; ++j){
                qout_M[v_size_*i+j] = 0;
                for(int k = 0; k < v_size_; ++k){
                    qout_M[v_size_*i+j] += m1[v_size_*i+k] * m2[v_size_*k + j];
                }
            }
        }
        dequantize(); // TODO complete dequantize function
    }
}
```

```
// 3) Call a function `blockMM()` to execute Matrix matrix multiplication
const float* ret = this->blockMM(comp);
```

```
struct Compute
{
    bool quantized;           // quantization or not
    int act_bits;             // precision of input values
    float act_min, act_max;    // min/max of input values
    int weight_bits;          // precision of weights
    float weight_min, weight_max; // min/max of weights
};
```

```
void quantize(float* input, char* quantized, int num_input, char bits_min, char bits_max,
char offset, float scale)
{
    for(int i = 0; i < num_input; i++)
    {
        quantized[i] = ?; // TODO: convert floating point to quantized value
    }
}

void dequantize(short* quantized, float* output, int num_output, char offset, float scale)
{
    for(int i = 0; i < num_output; i++)
    {
        output[i] = ?; // TODO: convert quantized value to floating point
    }
}
```

- [Note] Fill other functions `convLowering`, etc...

# Implement Quantization - FPGA

- Edit `src/fpga\_api.cpp`

```
const int *__attribute__((optimize("O0"))) FPGA::qblockMM(Compute* comp)
{
    num_block_call_ += 1;

    // fpga version
    *output_ = 0x5555;
    while (*output_ == 0x5555)
        ;

    return qdata_;
}
```

```
// 3) Call a function `blockMM()` to execute Matrix matrix multiplication
const int* ret = this->qblockMM(comp);
```

```
struct Compute
{
    bool quantized;           // quantization or not
    int act_bits;             // precision of input values
    float act_min, act_max;    // min/max of input values
    int weight_bits;          // precision of weights
    float weight_min, weight_max; // min/max of weights
};
```

```
void quantize(float* input, char* quantized, int num_input, char bits_min, char bits_max,
char offset, float scale)
{
    for(int i = 0; i < num_input; i++)
    {
        quantized[i] = ?; // TODO: convert floating point to quantized value
    }
}

void dequantize(short* quantized, float* output, int num_output, char offset, float scale)
{
    for(int i = 0; i < num_output; i++)
    {
        output[i] = ?; // TODO: convert quantized value to floating point
    }
}
```

- [Note] Fill other functions `convLowering` etc...

# Example: Download datasets

---

```
zed@debian-zynq:~/lab14$ bash download.sh
converted 'https://dl.dropbox.com/s/mdwy0kzf57nfl5f/t10k-images.idx3-ubyte' (ANSI_X3.4-1968) -> 'ht
--2019-06-03 22:17:32-- https://dl.dropbox.com/s/mdwy0kzf57nfl5f/t10k-images.idx3-ubyte
Resolving dl.dropbox.com (dl.dropbox.com)... 162.125.80.6, 2620:100:6030:6::a27d:5006
Connecting to dl.dropbox.com (dl.dropbox.com)|162.125.80.6|:443... connected.
HTTP request sent, awaiting response... 302 FOUND
Location: https://dl.dropboxusercontent.com/s/mdwy0kzf57nfl5f/t10k-images.idx3-ubyte [following]
converted 'https://dl.dropboxusercontent.com/s/mdwy0kzf57nfl5f/t10k-images.idx3-ubyte' (ANSI_X3.4-1
--2019-06-03 22:17:32-- https://dl.dropboxusercontent.com/s/mdwy0kzf57nfl5f/t10k-images.idx3-ubyte
Resolving dl.dropboxusercontent.com (dl.dropboxusercontent.com)... 162.125.80.6, 2620:100:6030:6::a
Connecting to dl.dropboxusercontent.com (dl.dropboxusercontent.com)|162.125.80.6|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 7840016 (7.5M) [application/octet-stream]
Saving to: 'data/t10k-images.idx3-ubyte'
```

# Example: Build your implementations

---

```
zed@debian-zynq:~/lab14$ make
g++ -fPIC -std=c++11 -I ./include -I./proto -o build/caffe_dnn.o -c ./src/caffe_dnn.cpp
g++ -fPIC -std=c++11 -I ./include -I./proto -o build/tf_dnn.o -c ./src/tf_dnn.cpp
g++ -fPIC -std=c++11 -I ./include -I./proto -o build/py_lib.o -c ./src/py_lib.cpp
g++ -fPIC -std=c++11 -I ./include -I./proto -o build/fpga_api_on_cpu.o -c src/fpga_api_on_cpu.cpp
g++ -shared -o build/libpylib_cpu.so build/py_lib.o build/caffe.pb.o build/caffe_dnn.o build/tf_dnn.o
g++ -fPIC -std=c++11 -I ./include -I./proto -o build/fpga_api.o -c src/fpga_api.cpp
g++ -shared -o build/libpylib_fpga.so build/py_lib.o build/caffe.pb.o build/caffe_dnn.o build/tf_dnn.o
```

# Skeleton Code

---

```
include
├── caffe_dnn.h
├── common_dnn.h
├── compute.h
├── fpga_api.h
├── ops.h
├── py_lib.h
└── tf_dnn.h
```

- Edit if you need -> fpga\_api.h

```
src
├── caffe_dnn.cpp
├── common_dnn.cpp
├── fpga_api.cpp
├── fpga_api_on_cpu.cpp
├── py_lib.cpp
└── tf_dnn.cpp
```

- Edit -> fpga\_api.cpp
- Edit -> fpga\_api\_cpu.cpp

# Checklists

---

- Convolution Lowering

- \$ sudo python eval.py --num\_test\_images 100 --m\_size 8 --v\_size 8 --  
**network cnn** --run\_type [cpu|fpga]

- **Quantization**

- \$ sudo python eval.py --num\_test\_images 100 --m\_size 8 --v\_size 8 --  
network cnn --run\_type [cpu|fpga] **--quantized**



# Example: Benchmark

- \$ sudo bash benchmark.sh

```
[*] Arguments: Namespace(a_bits=8, m_size=8, network='cnn', num_test_images=100, quantized=True, run_type='cpu', v_size=8, w_bits=8)
[*] Read MNIST...
[*] The shape of image: (100, 28, 28)
[*] Load the network...
[*] Run tests...
[*] Statistics...
{'accuracy': 1.0,
 'avg_num_call': 2206,
 'm_size': 8,
 'total_image': 100,
 'total_time': ██████████
 'v_size': 8}
```

```
[*] Arguments: Namespace(a_bits=8, m_size=8, network='cnn', num_test_images=100, quantized=True, run_type='fpga', v_size=8, w_bits=8)
[*] Read MNIST...
[*] The shape of image: (100, 28, 28)
[*] Load the network...
[*] Run tests...
[*] Statistics...
{'accuracy': 1.0,
 'avg_num_call': 2206,
 'm_size': 8,
 'total_image': 100,
 'total_time': ██████████
 'v_size': 8}
-e
=> Accuracy should be 1.0
```

# Optional Project - Quantization

This is an optional project that you can challenge if you are interested in improving performance

# Optimize your MM

---

- Enable quantization working on CPU / FPGA
  - Quantization on CPU: Implement on `fpga_api_cpu.cpp`
  - Quantization on FPGA: Implement on `fpga_api.cpp`
- Verilog code -> Int8 MM
  - Your data is quantized into 8bit, and it is transferred to your bitstream.
  - You don't need to use IP catalog, just implement the Int 8bit MM.
- If you are curious or confused about something,  
Just question about it!